

Introduction to Parallel Programming with NVIDIA CUDA

NVIDIA CUDA

License / Attribution




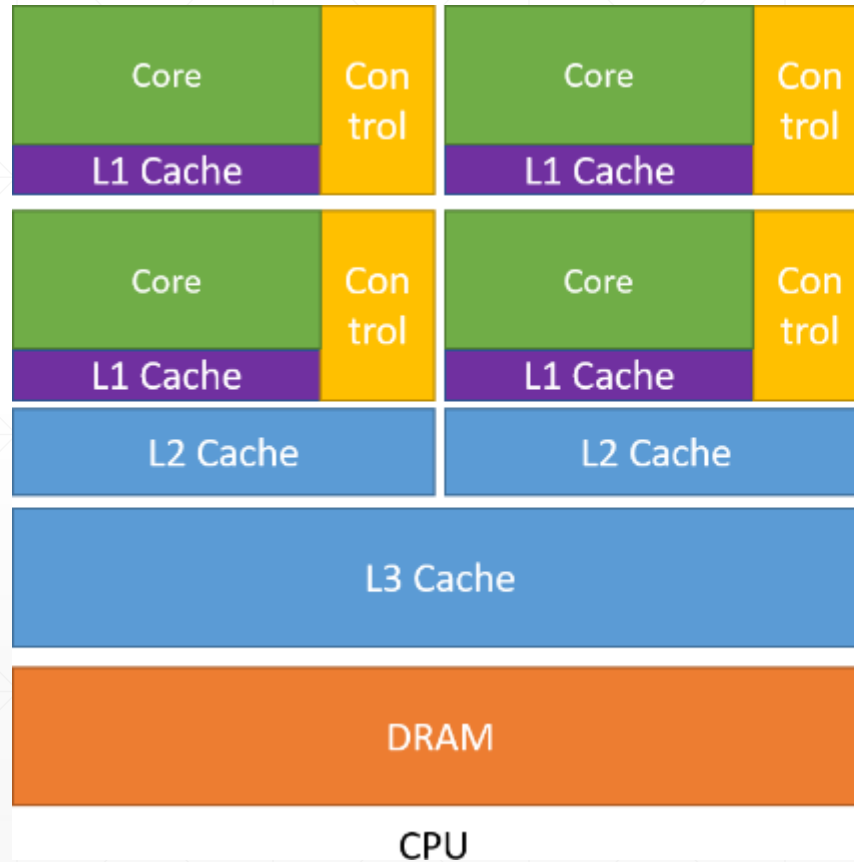
- Materials for the short-course „**Digital Signal Processing with GPUs — Introduction to Parallel Programming**” are licensed by us4us Ltd. the IPPT PAN under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).
- Some slides and examples are borrowed from the course „**The GPU Teaching Kit**” that is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).
 - All the borrowed slides are marked with  

GPU Architecture

Flynn Taxonomy of parallelism

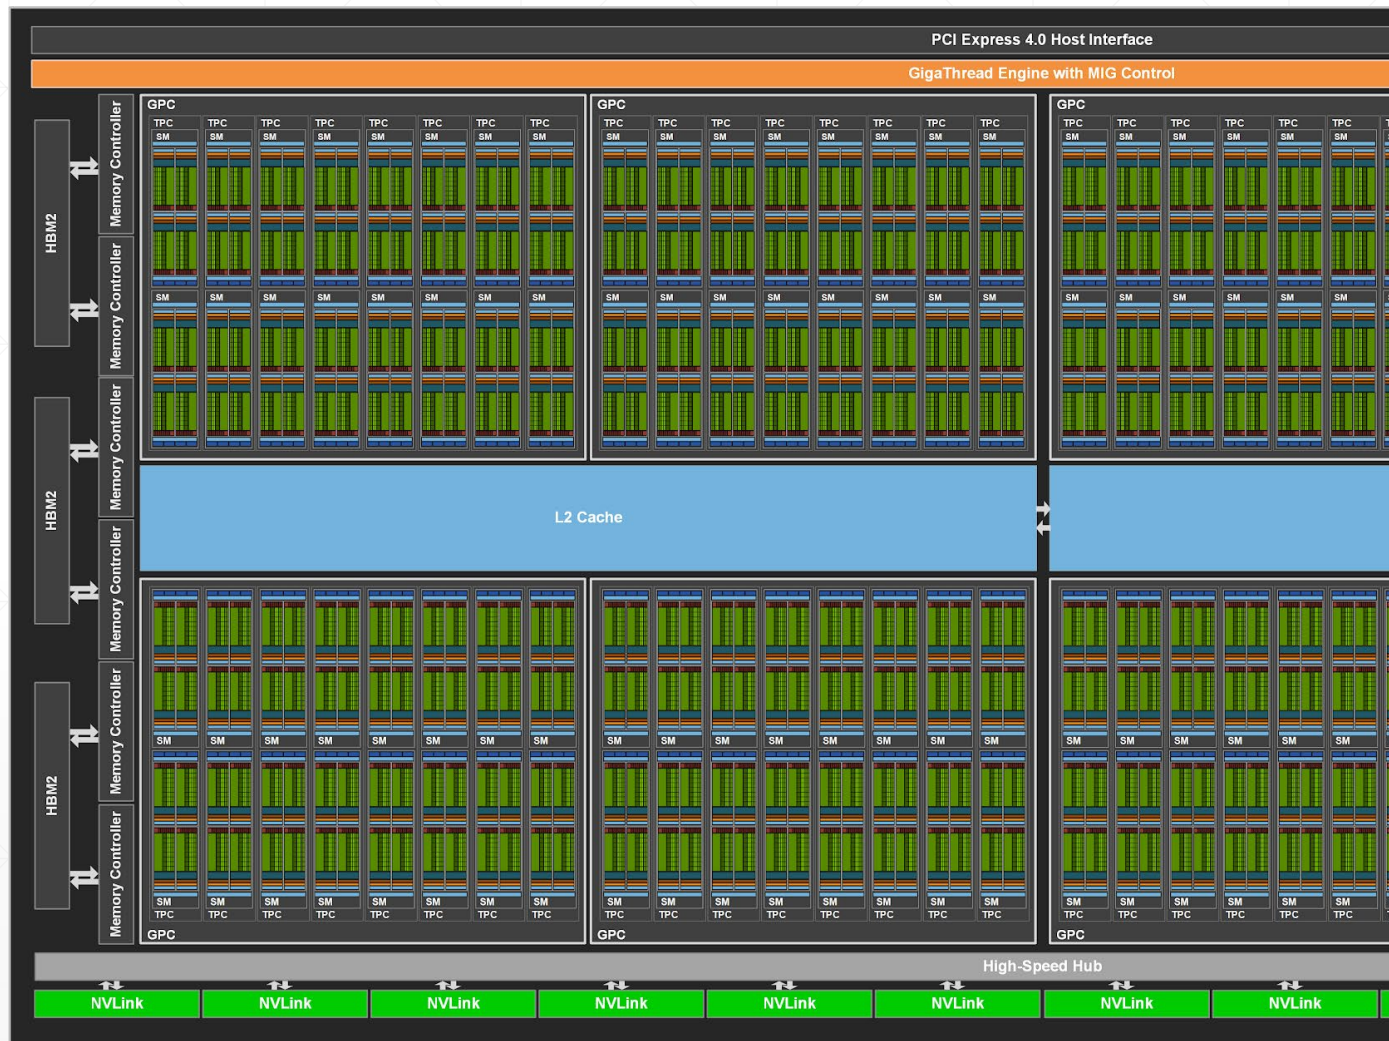
- Two dimensions
 - Number of instruction streams: single vs. multiple
 - Number of data streams: single vs. multiple
- **SISD** – single-instruction single-data
 - Pipelining and ILP (Instruction Level Parallelism) on a uniprocessor
- **SIMD** – single-instruction multiple-data (aka Vector processor)
 - DLP (Data Level Parallelism) on a vector processor
- **MIMD** – multiple-instruction multiple-data
 - DLP, TLP (Thread Level Parallelism) on a parallel processor
 - SPMD: single-program multiple data

Architecture CPU vs. GPU



Source: <https://docs.nvidia.com/cuda/>

NVIDIA Ampere Architecture



Source: <https://developer.nvidia.com/blog/nvidia-ampere-architecture-in-depth/>

Just a few numbers ...

Data Center GPU	NVIDIA Tesla P100	NVIDIA Tesla V100	NVIDIA A100
GPU Codename	GP100	GV100	GA100
GPU Architecture	NVIDIA Pascal	NVIDIA Volta	NVIDIA Ampere
GPU Board Form Factor	SXM	SXM2	SXM4
SMs	56	80	108
TPCs	28	40	54
FP32 Cores / SM	64	64	64
FP32 Cores / GPU	3584	5120	6912
FP64 Cores / SM	32	32	32
FP64 Cores / GPU	1792	2560	3456
INT32 Cores / SM	NA	64	64
INT32 Cores / GPU	NA	5120	6912
Tensor Cores / SM	NA	8	4 ²
Tensor Cores / GPU	NA	640	432

Data center GPU	NVIDIA Tesla P100	NVIDIA Tesla V100	NVIDIA A100
GPU Codename	GP100	GV100	GA100
GPU Architecture	NVIDIA Pascal	NVIDIA Volta	NVIDIA Ampere
Compute Capability	6.0	7.0	8.0
Threads / Warp	32	32	32
Max Warps / SM	64	64	64
Max Threads / SM	2048	2048	2048
Max Thread Blocks / SM	32	32	32
Max 32-bit Registers / SM	65536	65536	65536
Max Registers / Block	65536	65536	65536
Max Registers / Thread	255	255	255
Max Thread Block Size	1024	1024	1024
FP32 Cores / SM	64	64	64
Ratio of SM Registers to FP32 Cores	1024	1024	1024
Shared Memory Size / SM	64 KB	Configurable up to 96 KB	Configurable up to 164 KB

Source: <https://developer.nvidia.com/blog/nvidia-ampere-architecture-in-depth/>


More on Scalability

- Performance growth with HW generations
 - Increasing number of compute units (cores)
 - Increasing number of threads
 - Increasing vector length
 - Increasing pipeline depth
 - Increasing DRAM burst size
 - Increasing number of DRAM channels
 - Increasing data movement latency

CUDA Architecture

What is CUDA® (Compute Unified Device Architecture)

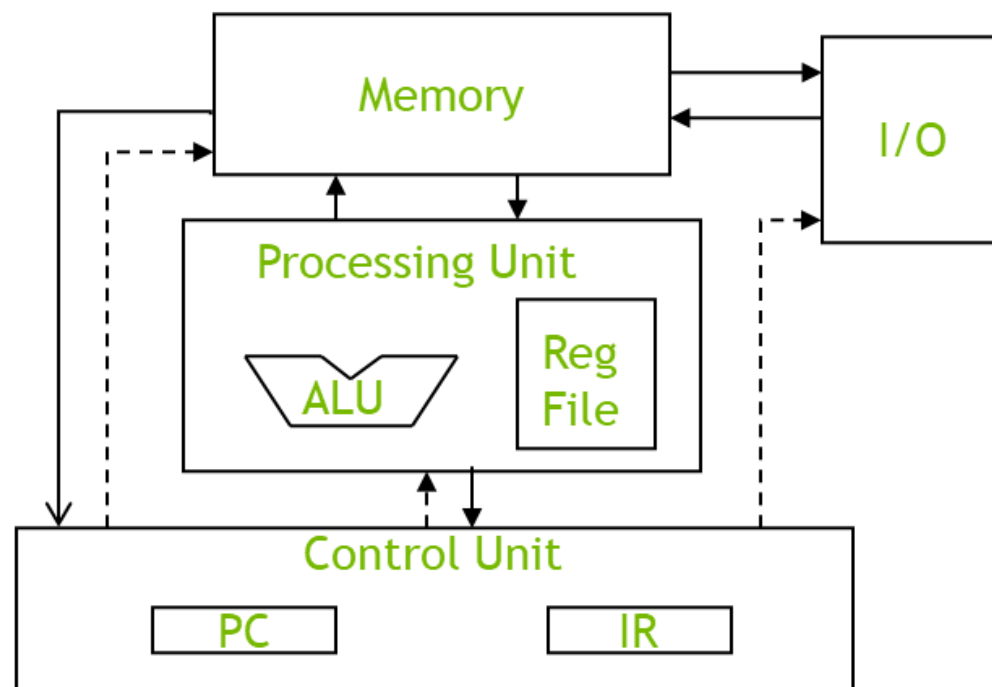
- NVIDIA CUDA is a General-Purpose Parallel Computing Platform and Programming Model
- Introduced back in 2006 by NVIDIA® to leverage the parallel compute engine in NVIDIA GPUs.
- CUDA is designed to support various languages and application programming interfaces.



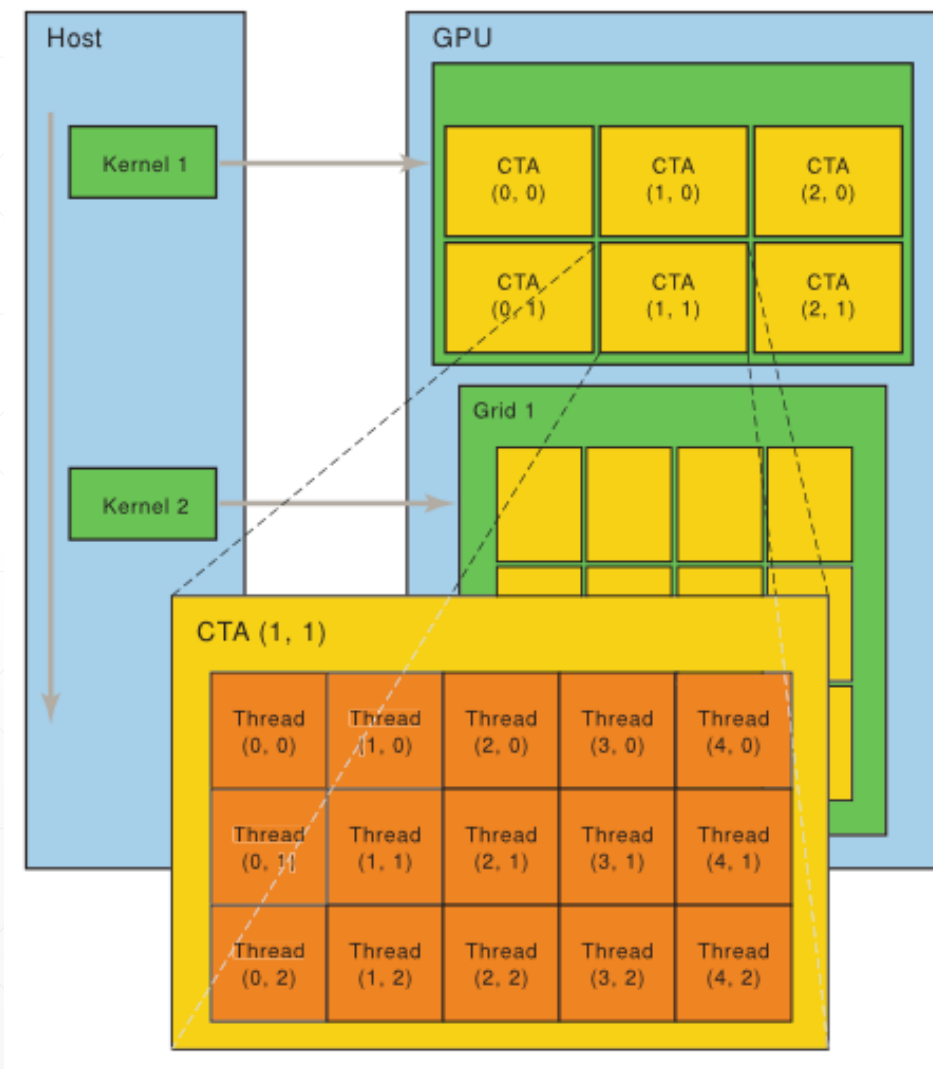
Source: <https://docs.nvidia.com/cuda/>

A Thread as a Von-Neumann Processor

A thread is a “virtualized” or
“abstracted”
Von-Neumann Processor





CUDA – Grid of Cooperative Thread Arrays



Source: <https://docs.nvidia.com/cuda/>

CUDA / GPU Automatic Scalability

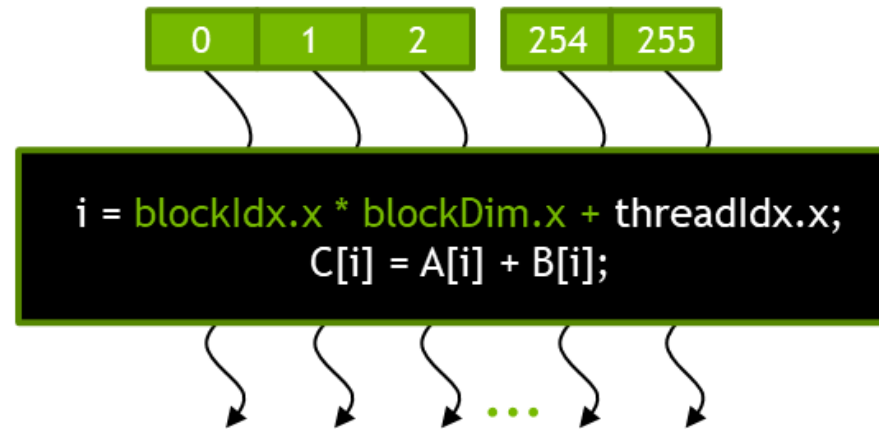
Grid of Thread Blocks



Source: <https://docs.nvidia.com/cuda/>


Arrays of Parallel Threads

- A CUDA kernel is executed by a **grid** (array) of threads
 - All threads in a grid run the same kernel code (Single Program Multiple Data)
 - Each thread has indexes that it uses to compute memory addresses and make control decisions

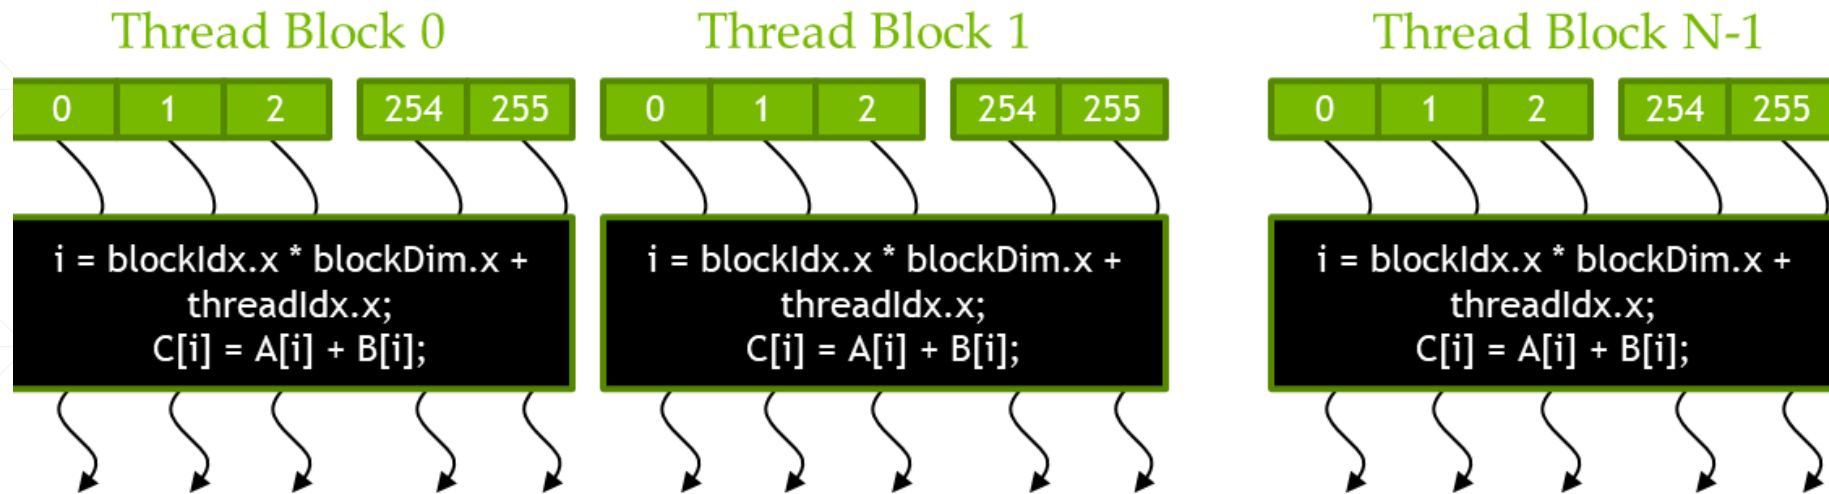


blockIdx and threadIdx

- Each thread uses indices to decide what data to work on
 - blockIdx: 1D, 2D, or 3D (CUDA 4.0)
 - threadIdx: 1D, 2D, or 3D
- Simplifies memory addressing when processing multidimensional data
 - Image processing
 - Solving PDEs on volumes
 - ...



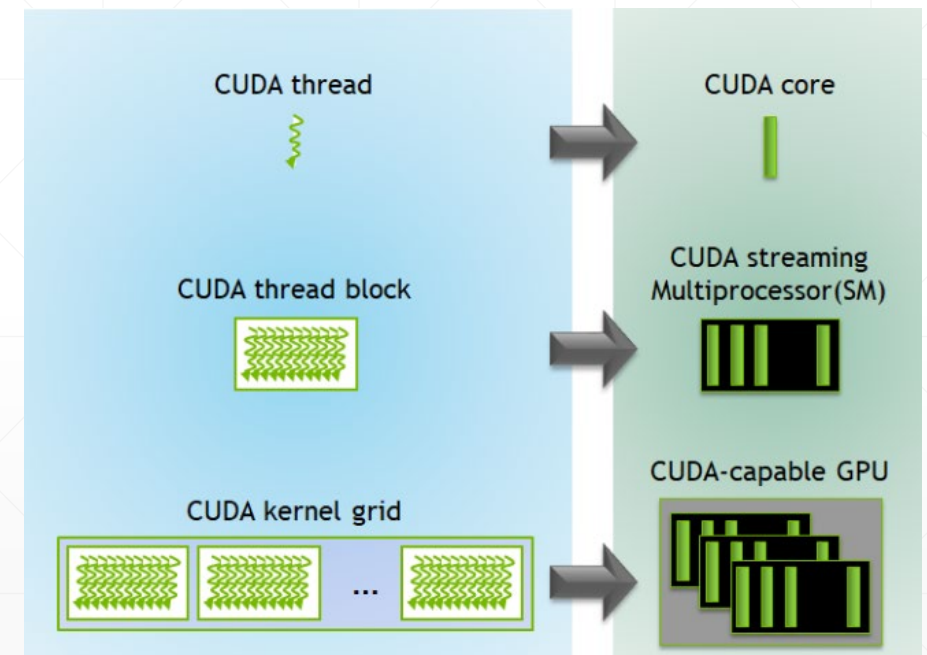
Thread Blocks: Scalable Cooperation



- Divide thread array into multiple blocks
 - Threads within a block cooperate via **shared memory, atomic operations** and **barrier synchronization**
 - Threads in different blocks do not interact

Thread Blocks / Grids

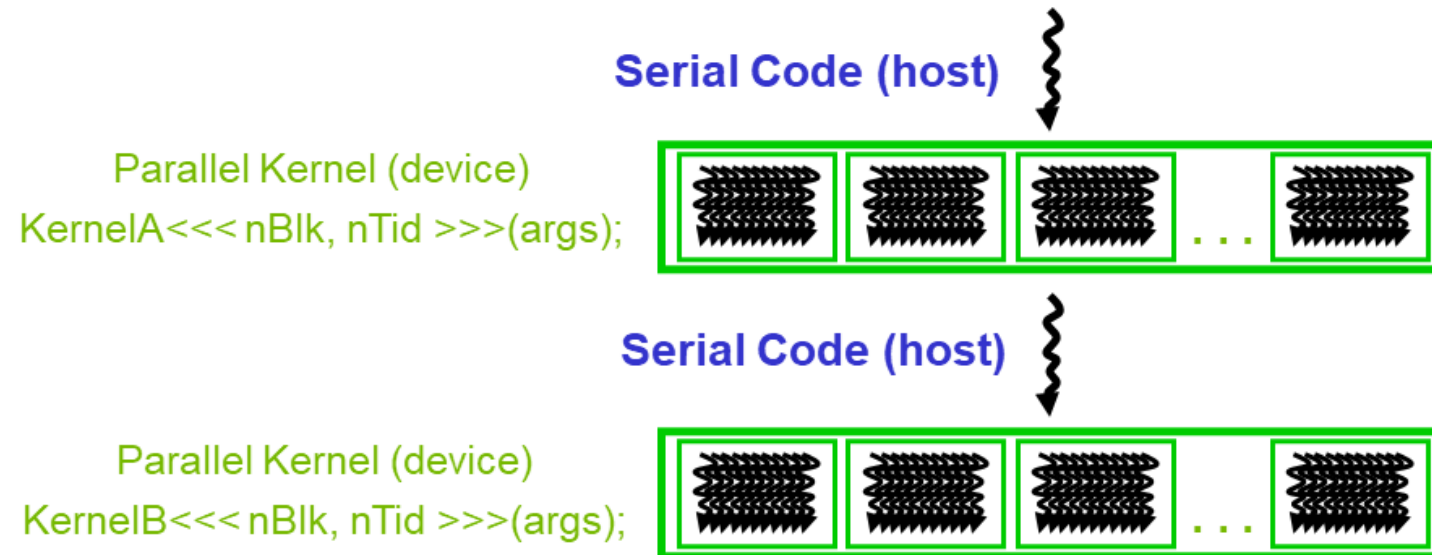
- Each CUDA block is executed by one streaming multiprocessor (SM) and cannot be migrated to other SMs in GPU.
- One SM can run several concurrent CUDA blocks depending on the resources needed by CUDA blocks.
- Each kernel is executed on one device and CUDA supports running multiple kernels on a device at one time.
- CUDA defines built-in 3D variables for threads and blocks
- CUDA architecture limits the numbers of threads per block (1024 threads per block limit).



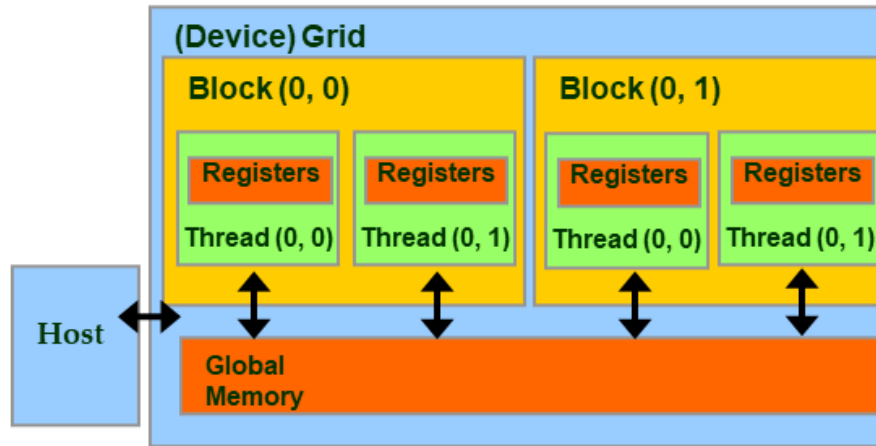
Source: <https://developer.nvidia.com/blog/cuda-refresher-cuda-programming-model/>

CUDA Execution Model

- Heterogeneous host (CPU) + device (GPU) application C program
 - Serial parts in **host** C code
 - Parallel parts in **device** SPMD kernel code



Partial Overview of CUDA Memories

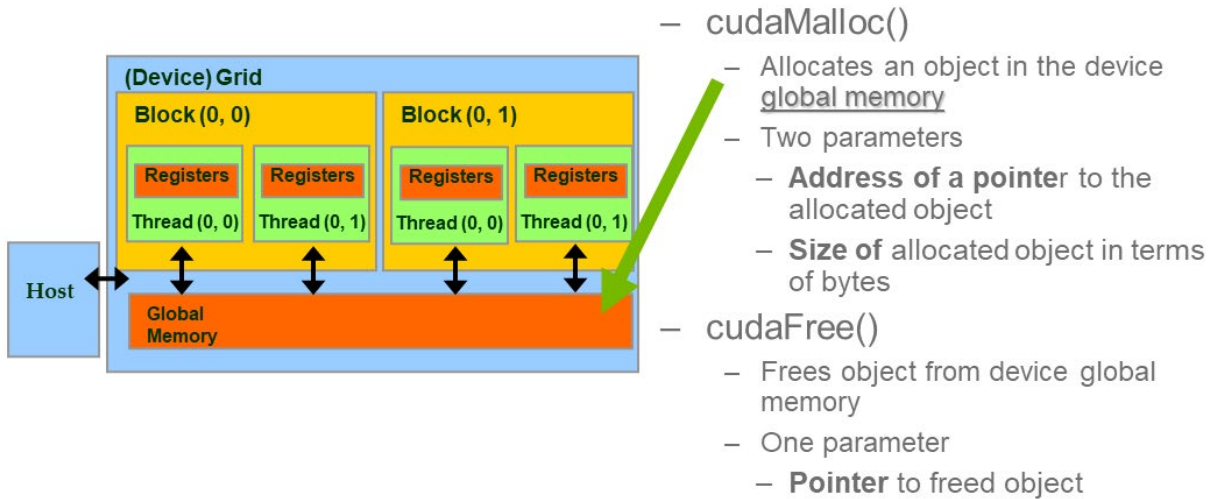


- Device code can:
 - R/W per-thread **registers**
 - R/W all-shared **global memory**
- Host code can
 - Transfer data to/from per grid **global memory**

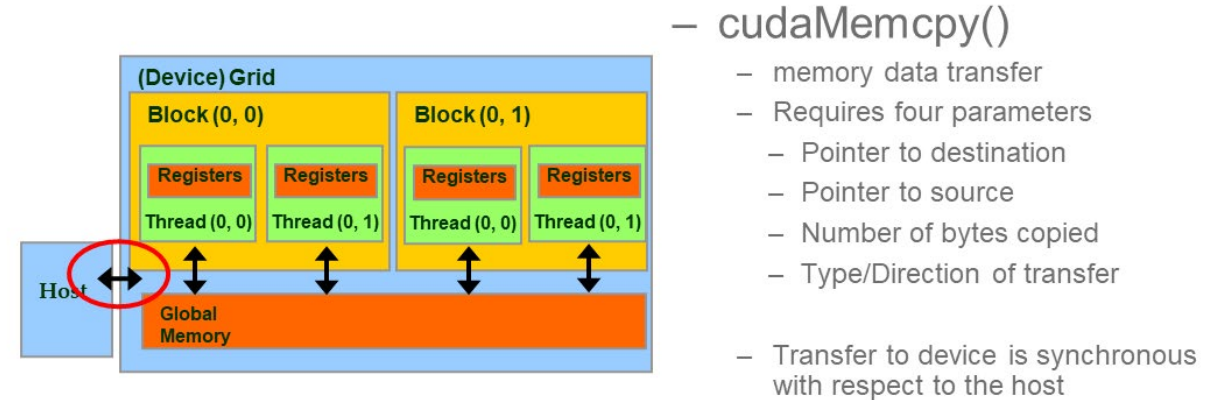
We will cover more memory types and more sophisticated memory models later.

Host/Device Memory


CUDA Device Memory Management API functions



Host-Device Data Transfer API functions



Data Parallelism – Vector Addition Example



Vector Addition – Traditional C Code

```
// Compute vector sum C = A + B
void vecAdd(float *h_A, float *h_B, float *h_C, int n)
{
    int i;
    for (i = 0; i < n; i++) h_C[i] = h_A[i] + h_B[i];
}

int main()
{
    // Memory allocation for h_A, h_B, and h_C
    // I/O to read h_A and h_B, N elements
    ...
    vecAdd(h_A, h_B, h_C, N);
}
```

Vector Addition, Explicit Memory Management

... Allocate h_A , h_B , h_C ...

```
void vecAdd(float *h_A, float *h_B, float *h_C, int n)
{
    int size = n * sizeof(float); float *d_A, *d_B, *d_C;
```

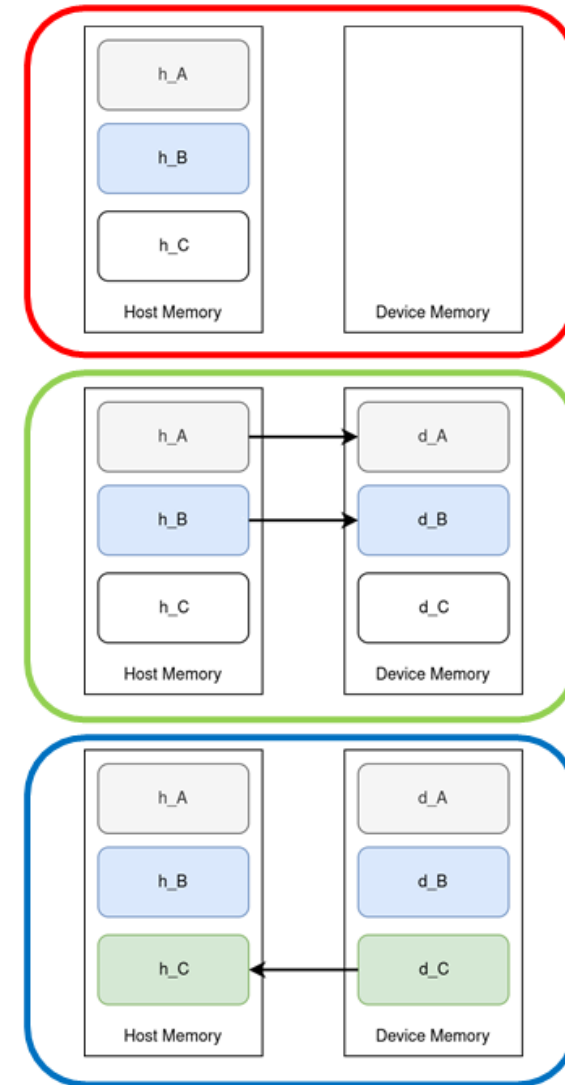
```
    cudaMalloc((void **) &d_A, size);
    cudaMalloc((void **) &d_B, size);
    cudaMalloc((void **) &d_C, size);
```

```
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
```


```
    // Kernel invocation code – to be shown later

    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);
    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
}
```

... Free h_A , h_B , h_C ...



Heterogeneous Computing vecAdd CUDA Host Code



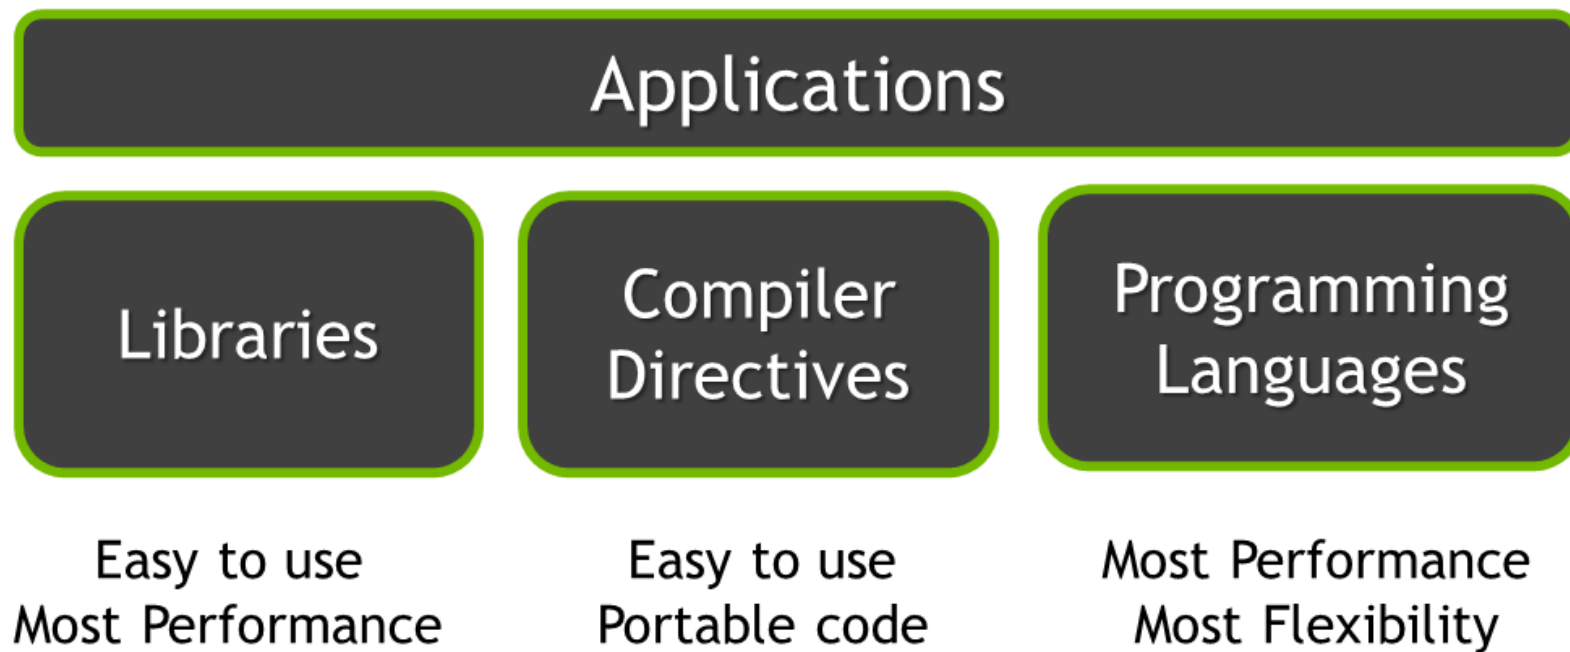
```
#include <cuda.h>
void vecAdd(float *h_A, float *h_B, float *h_C, int n)
{
    int size = n* sizeof(float);
    float *d_A, *d_B, *d_C;
    // Part 1
    // Allocate device memory for A, B, and C
    // copy A and B to device memory

    // Part 2
    // Kernel launch code – the device performs the actual
    vector addition

    // Part 3
    // copy C from the device memory
    // Free device vectors
}
```


GPU Acceleration

3 Ways to Accelerate Applications



NVIDIA GPU Accelerated Libraries

DEEP LEARNING





LINEAR ALGEBRA




SIGNAL, IMAGE, VIDEO



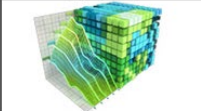
PARALLEL ALGORITHMS



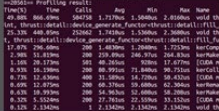
Developer Tools - Profilers



NSIGHT




NVVP




NVPROF

NVIDIA Provided



TAU

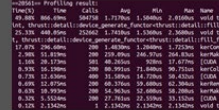


VampirTrace

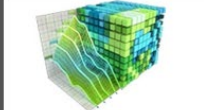
3rd Party

<https://developer.nvidia.com/performance-analysis-tools>

Profiling Tools




nvprof




NVVP

Phasing out



Nsight Compute




Nsight Systems

Current

See lecture 2-4 for an overview of all tools

Optimization



GPU Programming Languages

Numerical analytics ▶

MATLAB, Mathematica, LabVIEW

Python ▶

PyCUDA, Numba

Fortran ▶

CUDA Fortran, OpenACC

C ▶

CUDA C, OpenACC

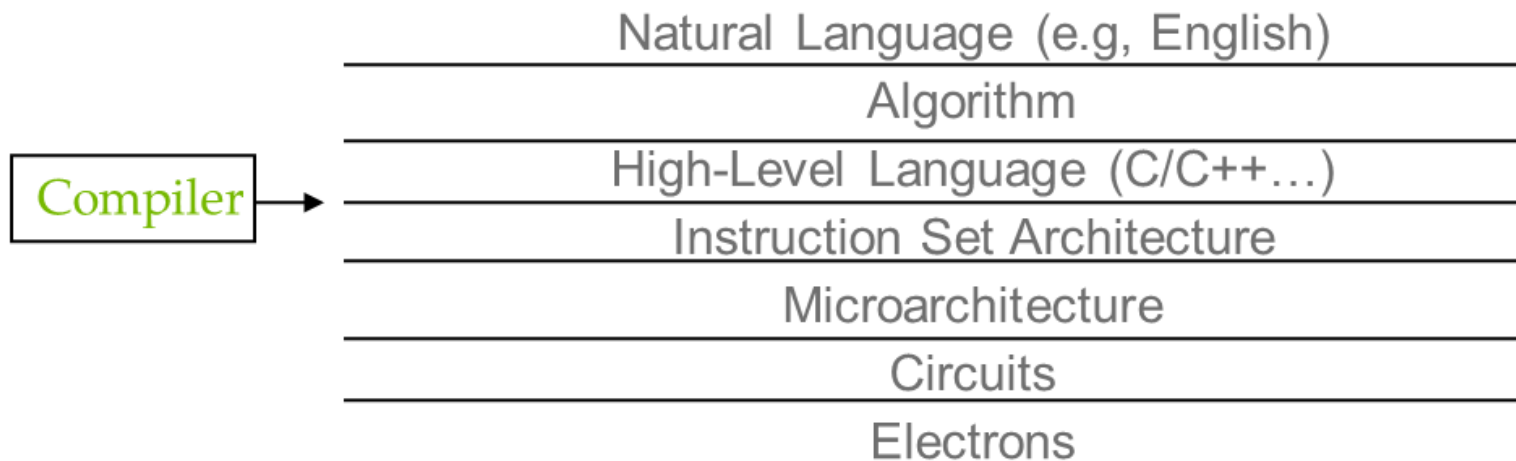
C++ ▶

CUDA C++, Thrust

C# ▶

Hybridizer

From Natural Language to Electrons



©Yale Patt and Sanjay Patel, *From bits and bytes to gates and beyond*

Why Python!?

Why is Python so Popular?

1. Easy to learn and use.
2. Mature and supportive Python Community.
3. Support from Big-Players (Corporate Sponsors).
4. “Batteries Included” – hundreds of libraries and frameworks available.
5. Versatility, efficiency, reliability, and speed. 🍀
6. Big data, machine learning and Cloud computing.
7. First-choice Language (see the ranking!).
8. The flexibility of Python language.
9. Use of Python in academics.
10. Automation

Language Ranking: IEEE Spectrum

2020

Rank	Language	Type	Score
1	Python▼	🌐 🖥️ ⚙️	100.0
2	Java▼	🌐 📱 🖥️	95.3
3	C▼	📱 🖥️ ⚙️	94.6
4	C++▼	📱 🖥️ ⚙️	87.0
5	JavaScript▼	🌐	79.5
6	R▼	🖥️	78.6
7	Arduino▼	⚙️	73.2
8	Go▼	🌐 🖥️	73.1
9	Swift▼	📱 🖥️	70.5
10	Matlab▼	🖥️	68.4

source: <https://spectrum.ieee.org/static/interactive-the-top-programming-languages-2020>

There should be one – and preferably only one – obvious way to do it (The Zen of Python) 😊

- There are many solutions for GPU acceleration in Python ...
- Python CUDA programming:
 - **NUMBA**, pyCUDA, pyOpenCL
- Libraries:
 - Numpy on the GPU: CuPy
 - Numpy on the GPU (again): Jax
 - Pandas on the GPU: RAPIDS cuDF
 - Scikit-Learn on the GPU: RAPIDS cuML
- Also deep learning frameworks like PyTorch, TensorFlow, Caffe, MXNet

NUMBA

- Accelerate Python Functions
- Numba translates Python functions to optimized machine code at runtime using the industry-standard [LLVM](#) compiler library. Numba-compiled numerical algorithms in Python can approach the speeds of C or FORTRAN.
- You don't need to replace the Python interpreter, run a separate compilation step, or even have a C/C++ compiler installed. Just apply one of the Numba decorators to your Python function, and Numba does the rest.



Numba makes Python code fast

Numba is an open source JIT compiler that translates a subset of Python and NumPy code into fast machine code.

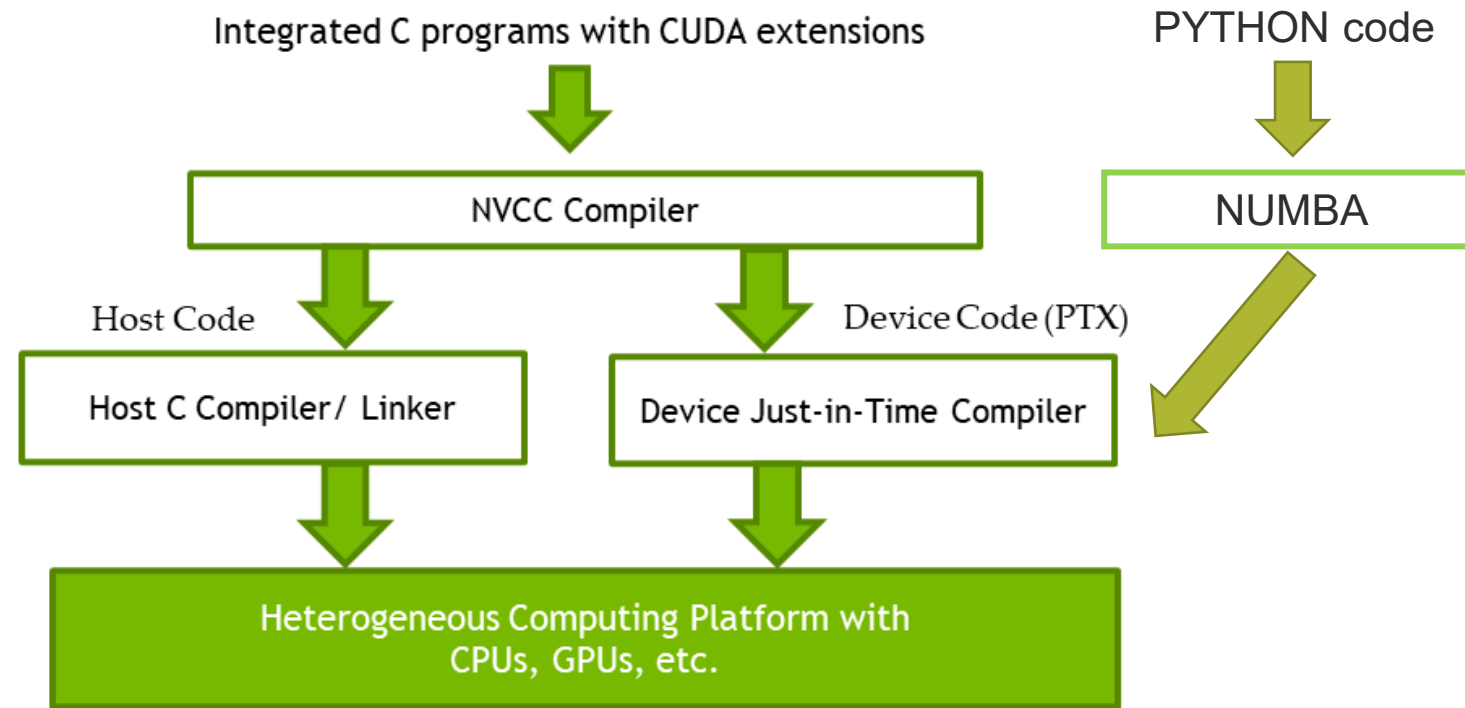
[Learn More](#)

[Try Numba »](#)

```
from numba import jit
import random

@jit(nopython=True)
def monte_carlo_pi(nsamples):
    acc = 0
    for i in range(nsamples):
        x = random.random()
        y = random.random()
        if (x ** 2 + y ** 2) < 1.0:
            acc += 1
    return 4.0 * acc / nsamples
```

Compiling A CUDA Program



GOOGLE Colab & Python

- Install Python + CUDA locally.
- Without installation, you can use Python in the CLOUD
 - For GOOGLE COLAB start here: <https://colab.research.google.com/notebooks/intro.ipynb>
- Other options in the Cloud: <https://developer.nvidia.com/how-to-cuda-python>
- Many other options and good developers' tools are available.